

# Deep Learning (CS324)


## 9. Generative adversarial networks<sup>\*</sup>

Prof. Jianguo Zhang  
SUSTech

<sup>\*</sup>Based on <http://www.deeplearningbook.org> chapter 20.10.4 + various sources cited in the slides

# Generative tasks

- Generation (from scratch): learn to sample from the distribution represented by the training set
  - *Unsupervised learning* task



# Generative tasks

- Conditional generation

goldfish



indigo  
bunting



redshank



saint  
bernard



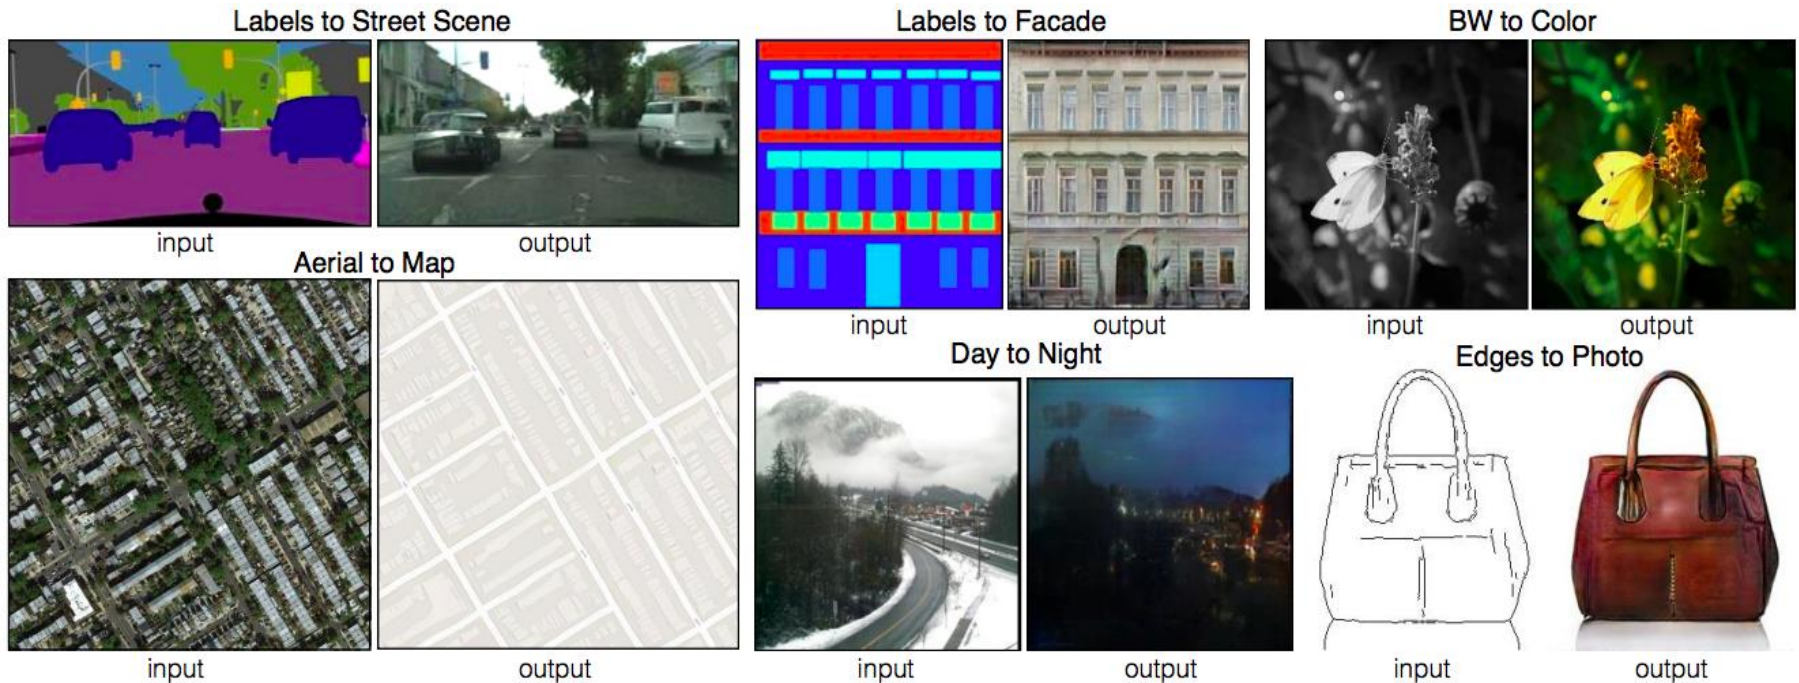
tiger  
cat






# Generative tasks

- Image-to-image translation



# Designing a network for generative tasks

- We saw that VAEs can learn to generate data by training both an encoder and a decoder (generator)
- Can we learn just the generator?
  - Recall the decoder of VAE



# Learning to sample



Training data  $x \sim p_{\text{data}}$

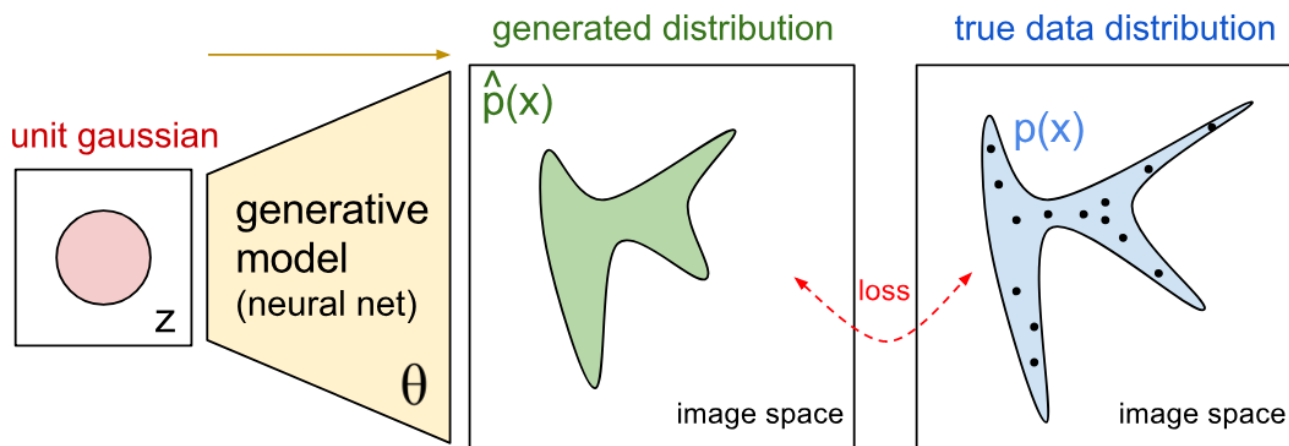


Generated samples  $x \sim p_{\text{model}}$

We want to learn  $p_{\text{model}}$  that matches  $p_{\text{data}}$

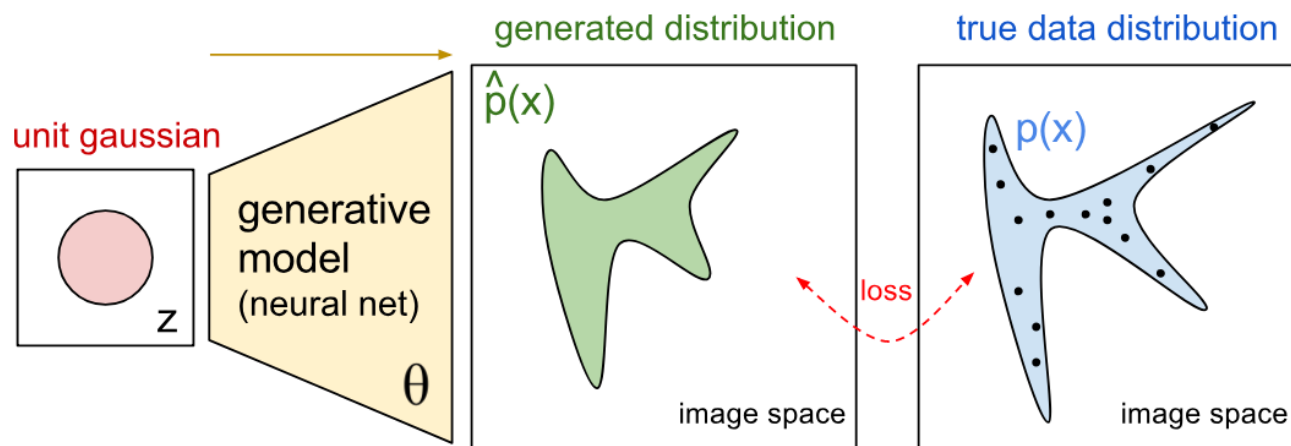
# From VAEs to GANs

- We saw that VAEs can learn to generate data by training both an encoder and a decoder (generator)
- Can we learn just the generator?



# From VAEs to GANs

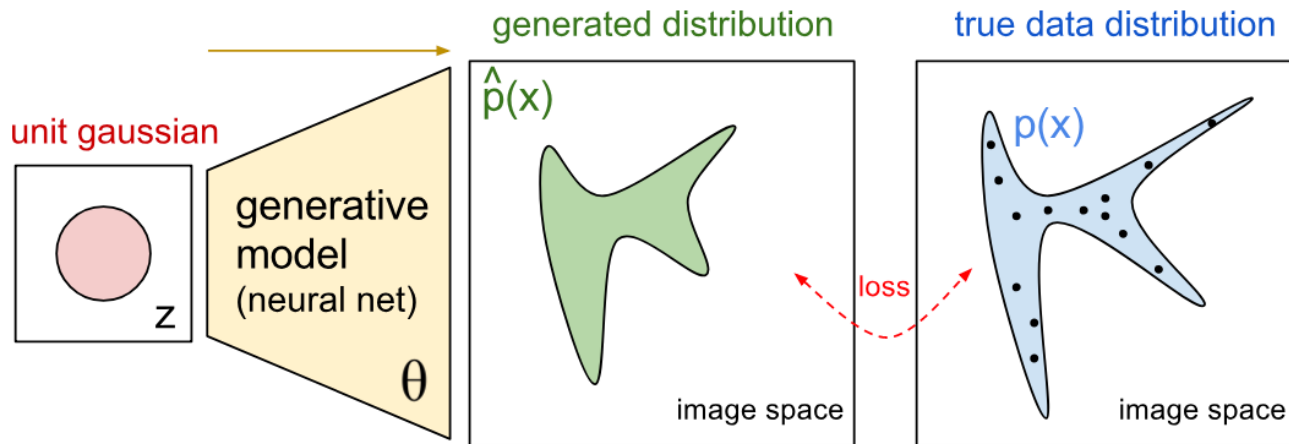
- But, how do we measure the quality of the generator without an encoder?
- In other words, what's the loss function?





# Generative adversarial networks

- GANs propose to learn the loss function
- The training process is a game between two networks



# Generative adversarial networks

- Train two networks with opposing objectives:
  - **Generator:** learns to generate samples
  - **Discriminator:** learns to distinguish between generated and real samples





Figure adapted from  
[F. Fleuret](#)

I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, Y. Bengio, [Generative adversarial nets](#), NIPS 2014


# Generative adversarial networks

- **Generator:**  $G(z)$  takes random noise  $z$  as input and outputs a (fake) image



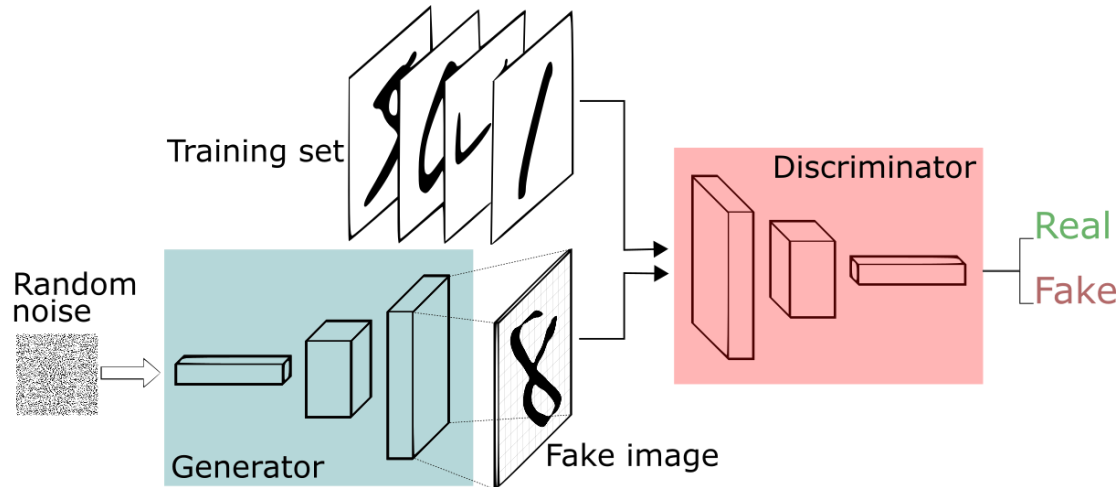
# Generative adversarial networks

- **Generator:**  $G(z)$  takes random noise  $z$  as input and outputs a (fake) image
- **Discriminator:**  $D(x)$  receives an image  $x$  in input, real or fake, and estimates its probability to be real




# Generative adversarial networks

- **Adversarial training:** the generator tries to fool the discriminator while the discriminator tries to get better at distinguishing fake vs real images



# Generative adversarial networks


- When the discriminator spots a fake the generator adjusts its parameters, until at the end the generator reproduces the true data distribution and the discriminator is unable to find differences






# Generative adversarial networks

- **Note:** both the generator and the discriminator need to be differentiable
- Typically both implemented as (deep) neural **networks**, so we can use backpropagation

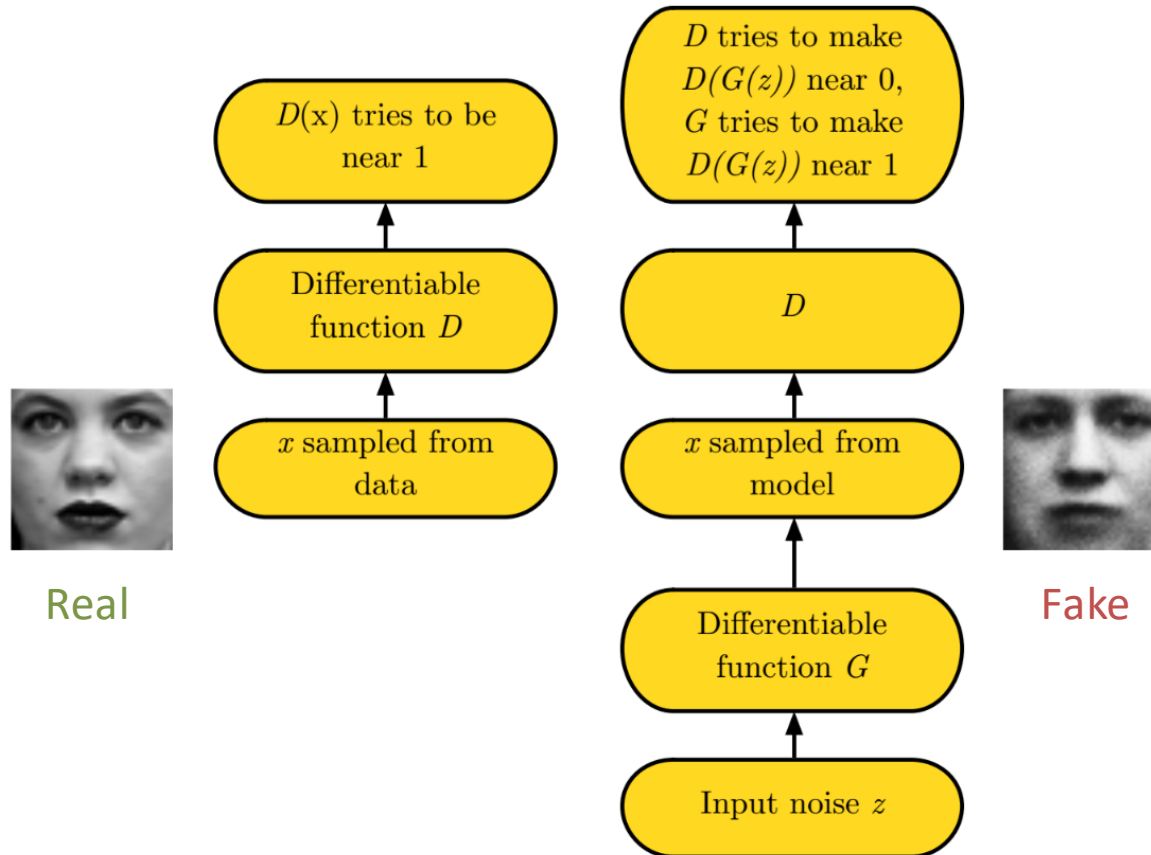



# Generative adversarial networks

- That's like saying that the discriminator at the end is unable to tell the difference between the generated data distribution and the true data distribution (where training data comes from)



# GANs pipeline







Real



Fake

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$





# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Prior distribution on input noise variables

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

The generator is a differentiable function (e.g., MLP) mapping noise to data space

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

The discriminator is a differentiable function (e.g., MLP) with single scalar output, i.e., the probability that  $\mathbf{x}$  comes from the true data distribution

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

$D(\mathbf{x}) = 1$  when the discriminator thinks  $\mathbf{x}$  is a real image

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

$D(G(\mathbf{z})) = 1$  when the discriminator thinks  $G(\mathbf{z})$  is a real image



# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

D is trained to maximise the probability of giving the correct label to samples from the true data distribution (the label should be 1) as well as to samples from G (the label should be 0)

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Large when  $D(\mathbf{x})$  close to 1

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Large when  $D(G(\mathbf{z}))$  close to 0

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Simultaneously G is trained to minimise  $\log(1 - D(G(\mathbf{z})))$ , i.e., fool the discriminator

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- For a fixed G, the loss is effectively the binary cross-entropy
- The minimax strategy is used for **zero-sum games**, i.e., loss of one player = gain of the adversary
- The minimax solution is the **Nash equilibrium**  
Nash equilibrium is a state in which D and G don't have any incentive to change because doing so would make the value of the objective function worse

# Minimax and zero-sum games

- G and D follow a **minimax strategy**

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- For a fixed G, the loss is effectively the binary cross-entropy
- The minimax strategy is used for **zero-sum games**, i.e., loss of one player = gain of the adversary
- The minimax solution is the **Nash equilibrium**  
In machine learning terms it means we reached convergence of the gradient descent for the joint optimisation problem



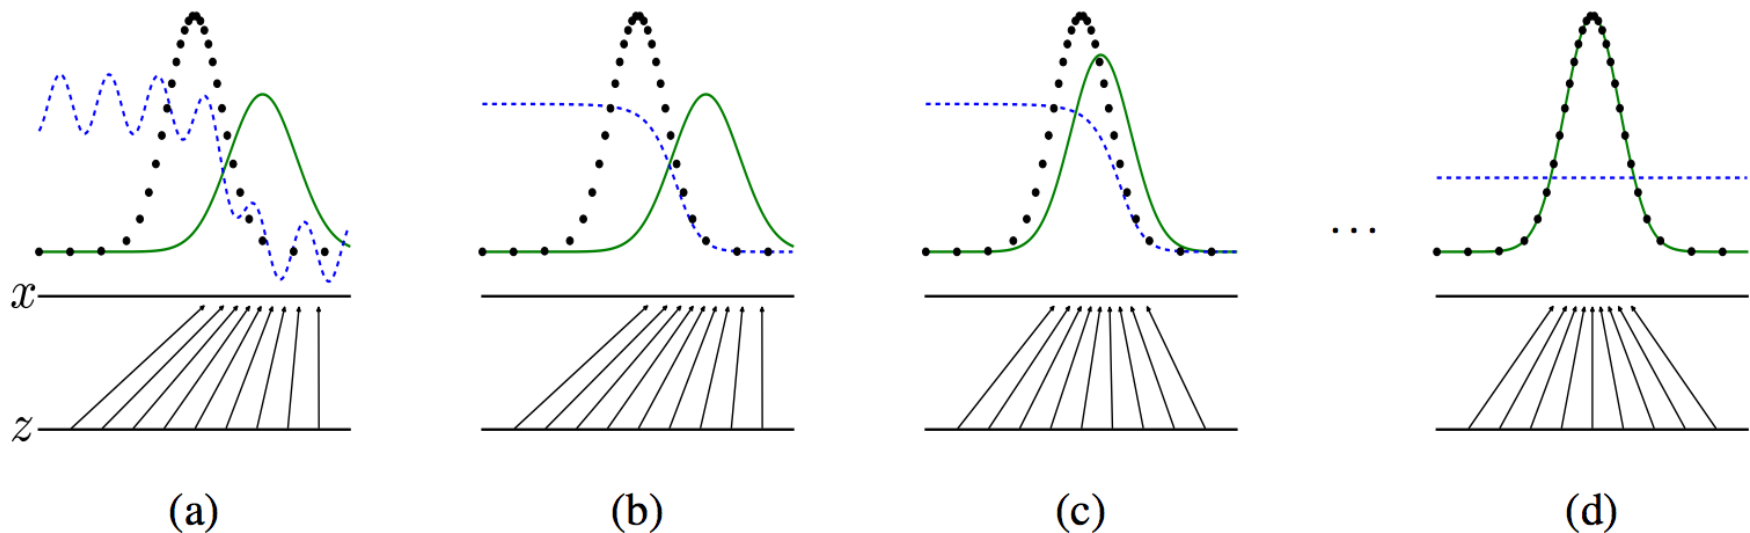


Figure 1: Generative adversarial nets are trained by simultaneously updating the discriminative distribution ( $D$ , blue, dashed line) so that it discriminates between samples from the data generating distribution (black, dotted line)  $p_{\mathbf{x}}$  from those of the generative distribution  $p_g$  ( $G$ ) (green, solid line). The lower horizontal line is the domain from which  $\mathbf{z}$  is sampled, in this case uniformly. The horizontal line above is part of the domain of  $\mathbf{x}$ . The upward arrows show how the mapping  $\mathbf{x} = G(\mathbf{z})$  imposes the non-uniform distribution  $p_g$  on transformed samples.  $G$  contracts in regions of high density and expands in regions of low density of  $p_g$ . (a) Consider an adversarial pair near convergence:  $p_g$  is similar to  $p_{\text{data}}$  and  $D$  is a partially accurate classifier. (b) In the inner loop of the algorithm  $D$  is trained to discriminate samples from data, converging to  $D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})}$ . (c) After an update to  $G$ , gradient of  $D$  has guided  $G(\mathbf{z})$  to flow to regions that are more likely to be classified as data. (d) After several steps of training, if  $G$  and  $D$  have enough capacity, they will reach a point at which both cannot improve because  $p_g = p_{\text{data}}$ . The discriminator is unable to differentiate between the two distributions, i.e.  $D(\mathbf{x}) = \frac{1}{2}$ .

# GAN training

- $V(G, D) = \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$

- Alternate between

- *Gradient ascent* on discriminator:

$$D^* = \arg \max_D V(G, D)$$

- *Gradient descent* on generator (minimize log-probability of discriminator being right):

$$\begin{aligned} G^* &= \arg \min_G V(G, D) \\ &= \arg \min_G \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \end{aligned}$$

- In practice, do *gradient ascent* on generator (maximize log-probability of discriminator being wrong):

$$G^* = \arg \max_G \mathbb{E}_{z \sim p} \log(D(G(z)))$$

# Training GANs

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# Training GANs

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log \left( 1 - D(G(z^{(i)})) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(z^{(i)})) \right).$$

In practice this saturates in the beginning since  $D$  rejects the samples with high confidence because they are clearly different from the training data

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# Training GANs

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (D(G(z^{(i)})))$$

So we use this loss function for the generator instead, i.e., the generator maximises the log-prob of the discriminator being mistaken


**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# GAN training

$$\min_{w_G} \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \text{ vs. } \max_{w_G} \mathbb{E}_{z \sim p} \log(D(G(z)))$$



# Example generator: DCGAN


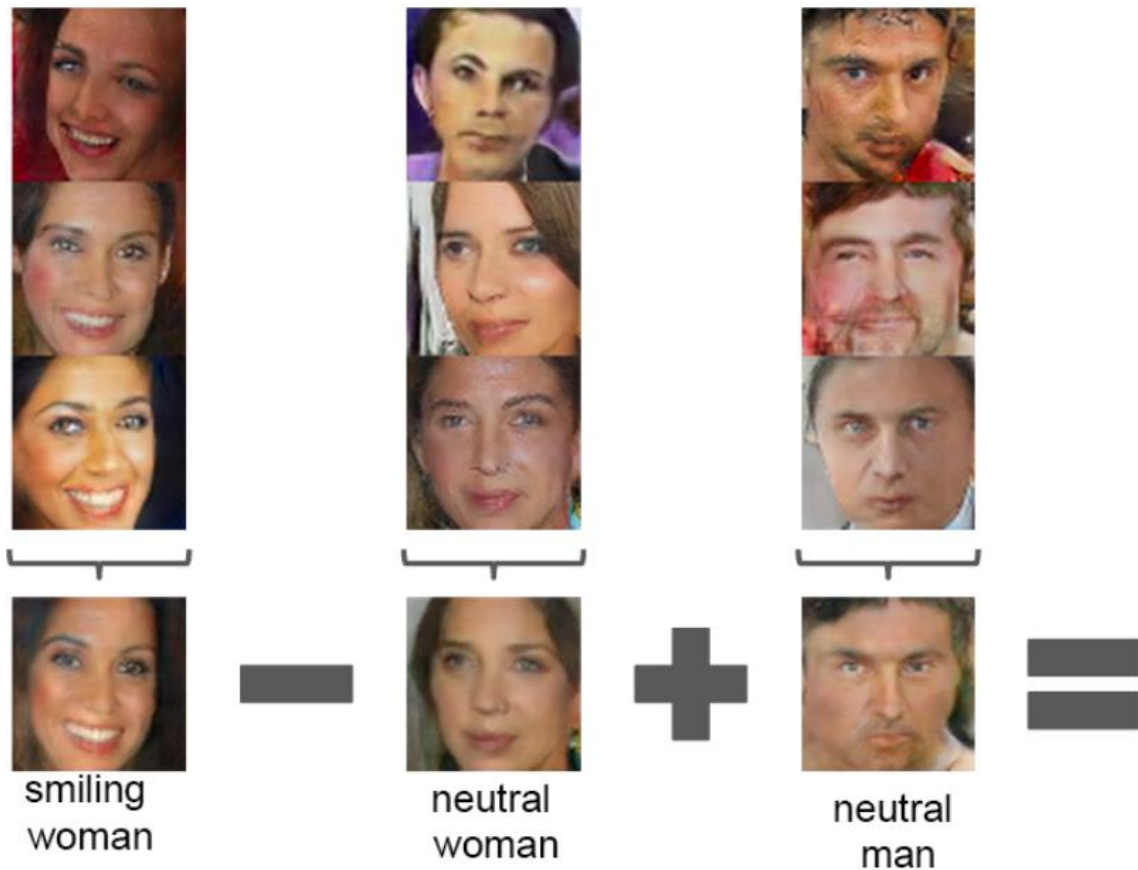


Figure 1: DCGAN generator used for LSUN scene modeling. A 100 dimensional uniform distribution  $Z$  is projected to a small spatial extent convolutional representation with many feature maps. A series of four fractionally-strided convolutions (in some recent papers, these are wrongly called deconvolutions) then convert this high level representation into a  $64 \times 64$  pixel image. Notably, no fully connected or pooling layers are used.

# DCGAN results

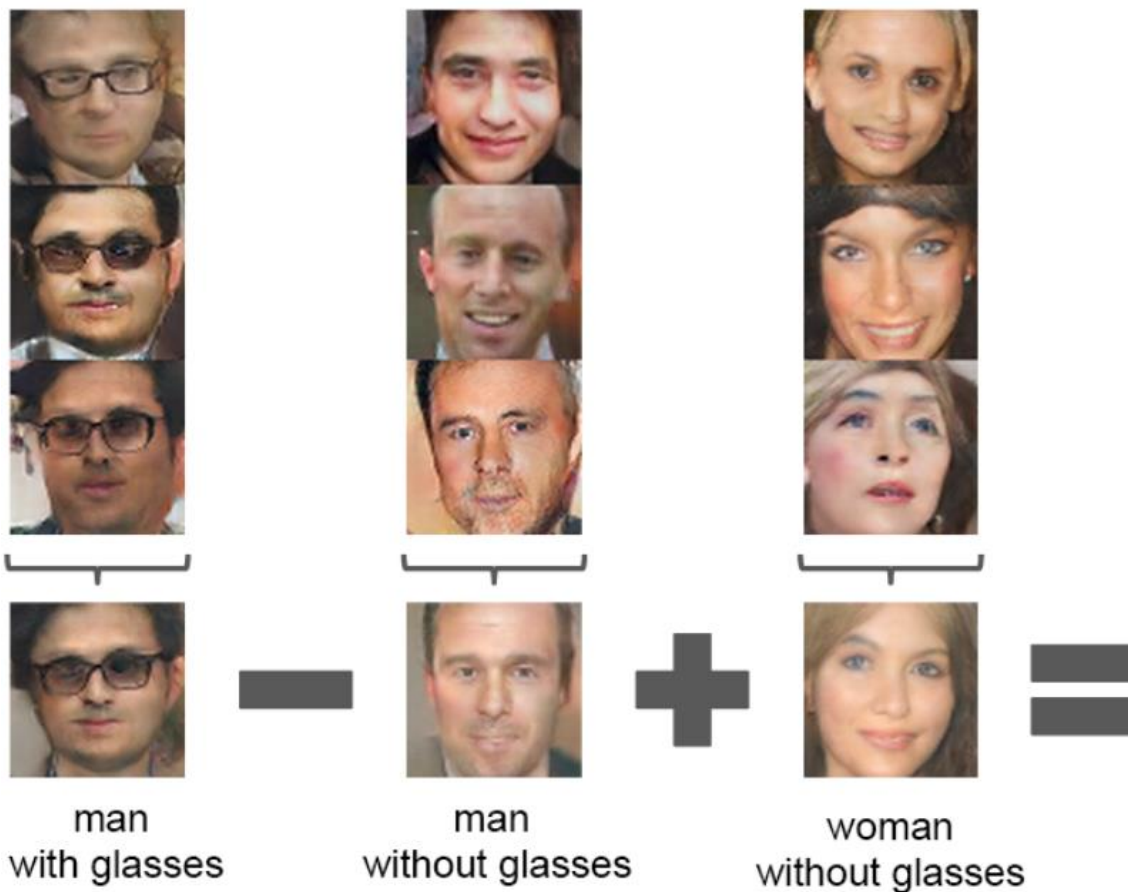
- Vector arithmetic in the z space





# DCGAN results

- Vector arithmetic in the z space



# DCGAN results

- Pose transformation by adding a “turn” vector



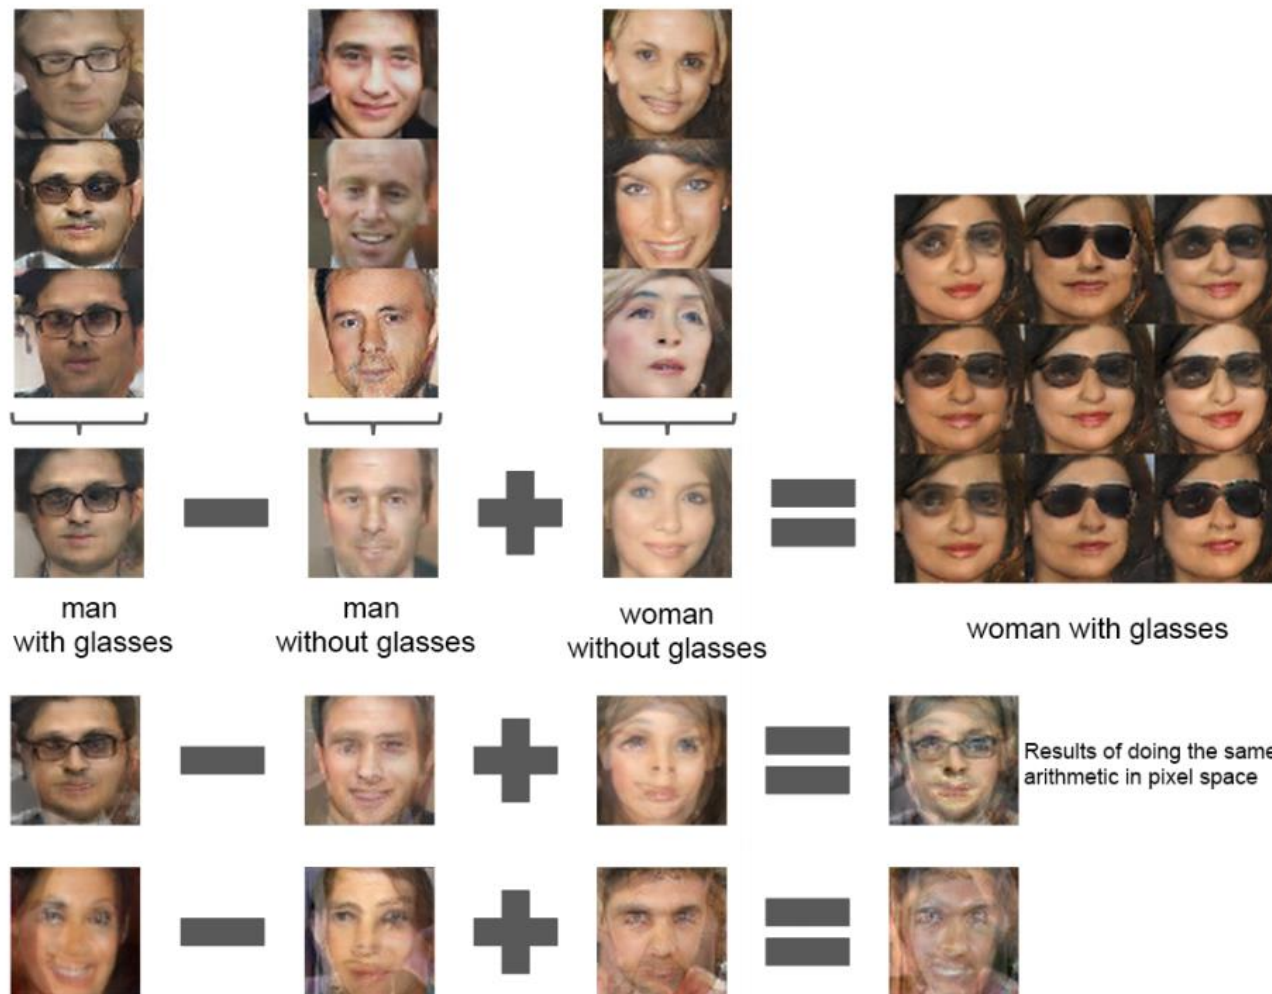


Figure 7: Vector arithmetic for visual concepts. For each column, the  $Z$  vectors of samples are averaged. Arithmetic was then performed on the mean vectors creating a new vector  $Y$ . The center sample on the right hand side is produce by feeding  $Y$  as input to the generator. To demonstrate the interpolation capabilities of the generator, uniform noise sampled with scale  $\pm 0.25$  was added to  $Y$  to produce the 8 other samples. Applying arithmetic in the input space (bottom two examples) results in noisy overlap due to misalignment.

# Problems with GANs

- In the real world we face multiple problems when training GANs:
  - **Finite sample size**: training set is finite, not the full distribution
  - **Limited capacity**: the generator has limit capacity, i.e. cannot perfectly represent any distribution
  - **Optimisation errors**: optimisers can get stuck in local optima or never exactly converge to global optima

# Problems with GANs

- In the real world we face multiple problems when training GANs:
  - **Saddle point problem**: harder than finding a maximum or minimum
  - **Balancing updates**: D too weak/strong means no gradient for G to improve

M. Arjovsky, S. Chintala, L. Bottou, [Wasserstein generative adversarial networks](#)



# Conditional generation (cGAN)

goldfish



indigo  
bunting



redshank



saint  
bernard




tiger  
cat




# Conditional generation (cGAN)

- Suppose we want to condition the generation of samples on discrete side information (label)  $y$ 
  - How do we add  $y$  to the basic GAN framework?



# Conditional generation

- Suppose we want to condition the generation of samples on discrete side information (label)  $y$ 
  - How do we add  $y$  to the basic GAN framework?





# Conditional generation

- Example: simple network for generating 28 x 28 MNIST digits

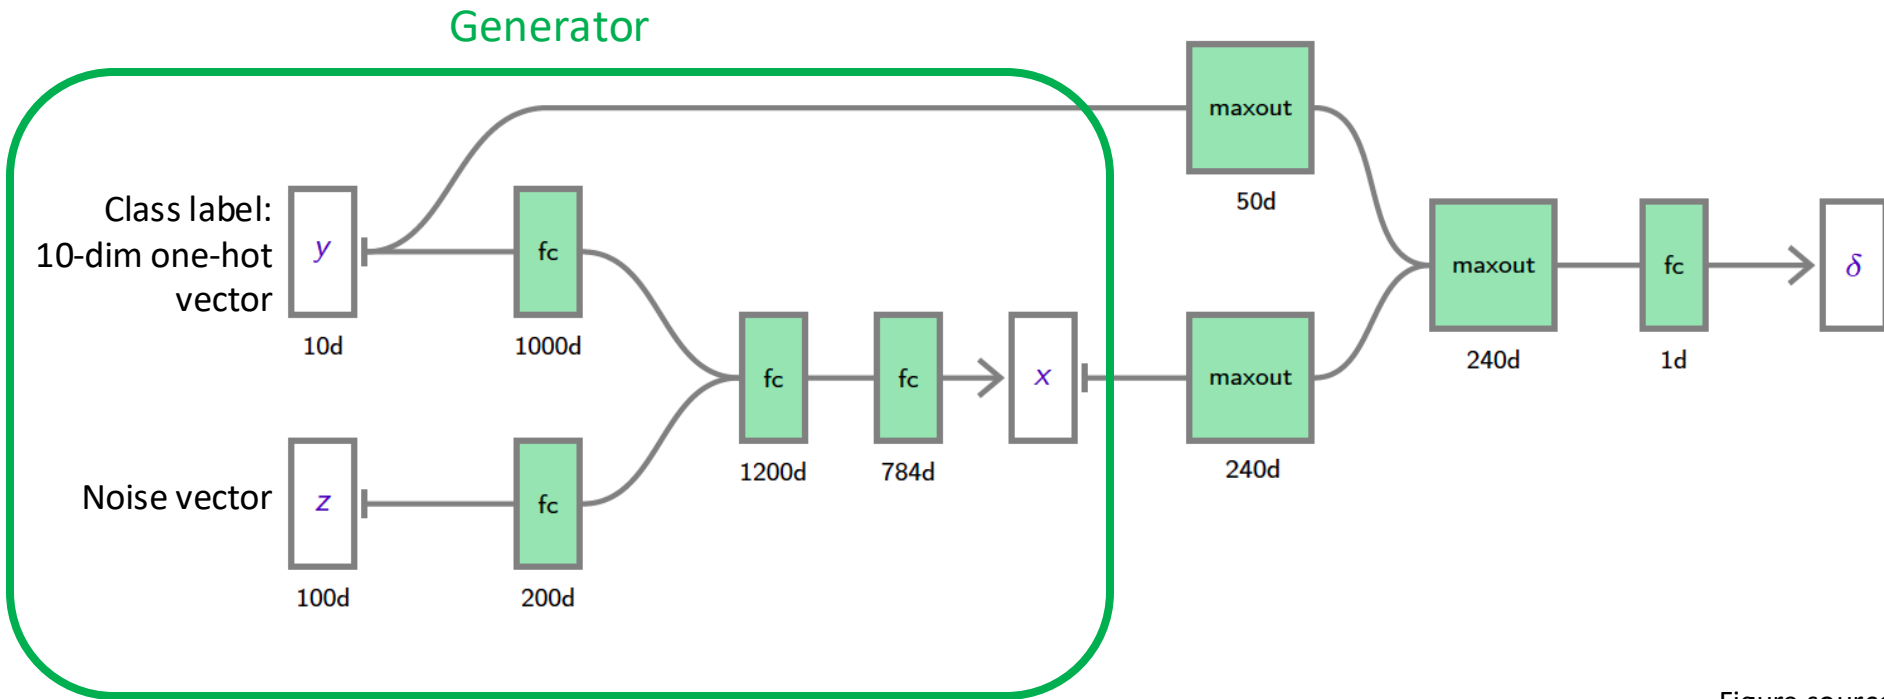


Figure source:  
[F. Fleuret](#)

# Conditional generation

- Example: simple network for generating 28 x 28 MNIST digits

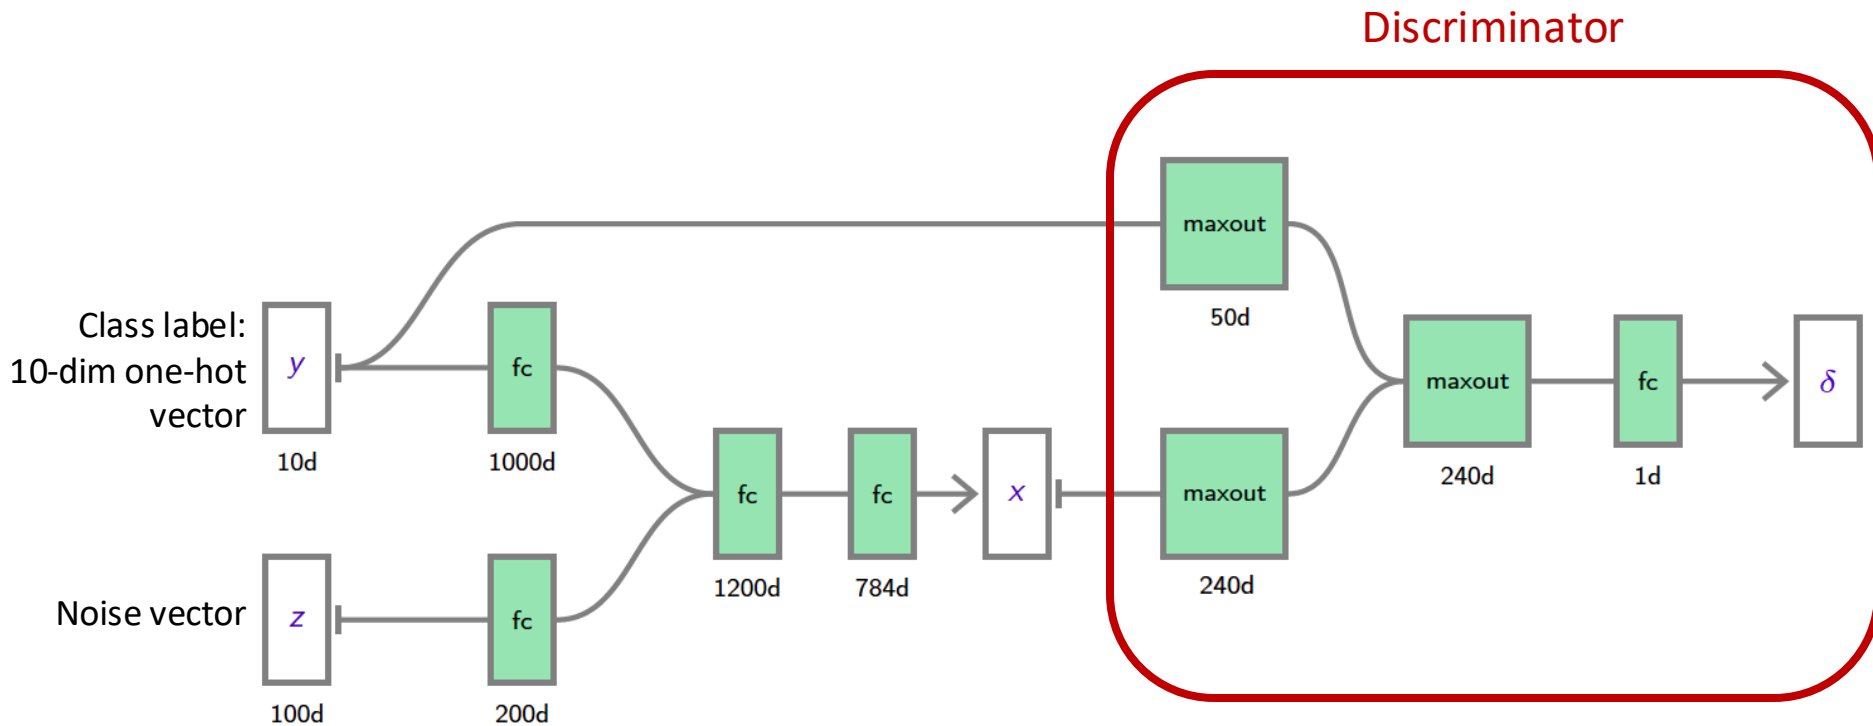


Figure source:  
[F. Fleuret](#)

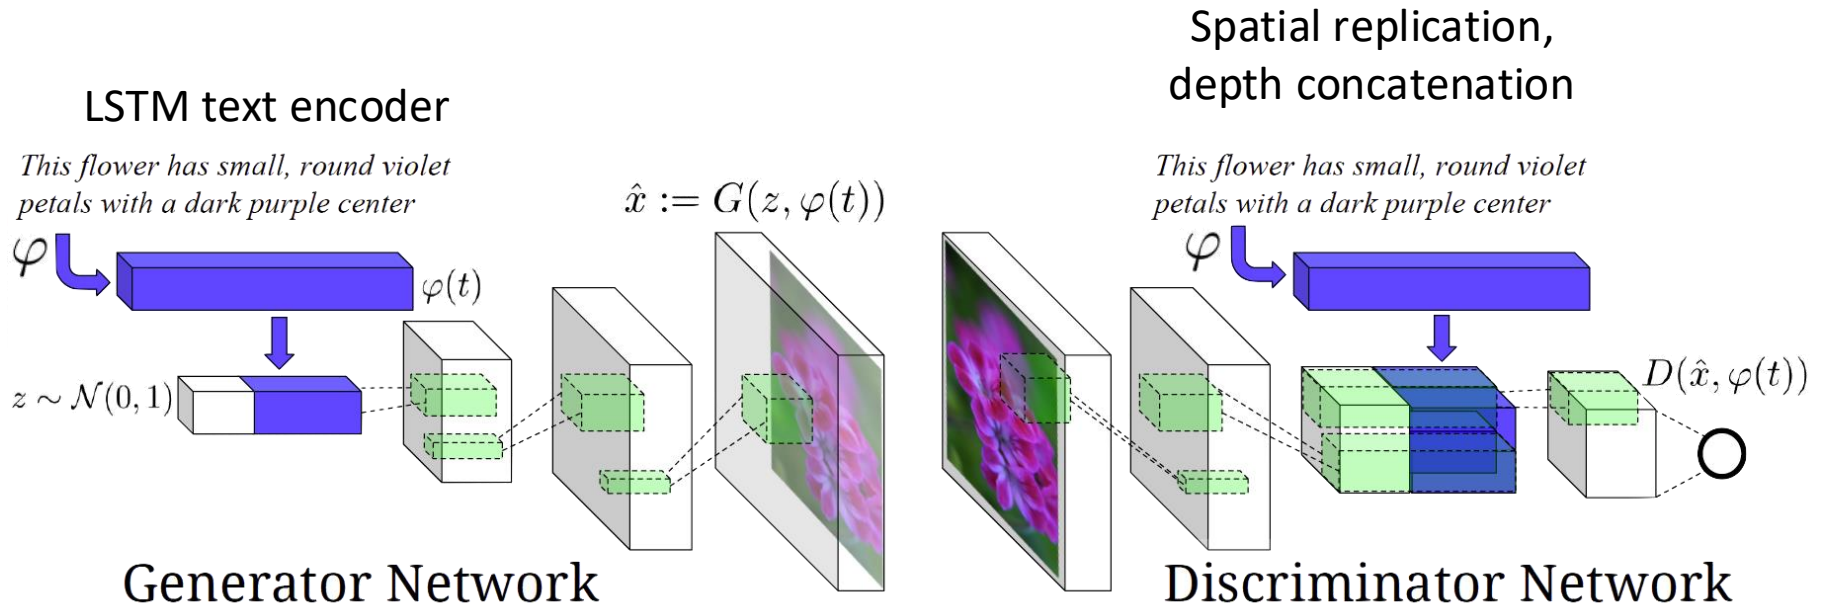
# Conditional generation

- Example: simple network for generating 28 x 28 MNIST digits



# Conditional generation

- Another example: text-to-image synthesis



# Conditional generation

- Another example: text-to-image synthesis

Previously unseen  
captions (*zero-shot*  
setting)

this small bird has a pink  
breast and crown, and black  
primaries and secondaries.



the flower has petals that  
are bright pinkish purple  
with white stigma



this magnificent fellow is  
almost all black with a red  
crest, and white cheek patch.




this white and yellow flower  
have thin white petals and a  
round yellow stamen



Captions seen in the  
training set

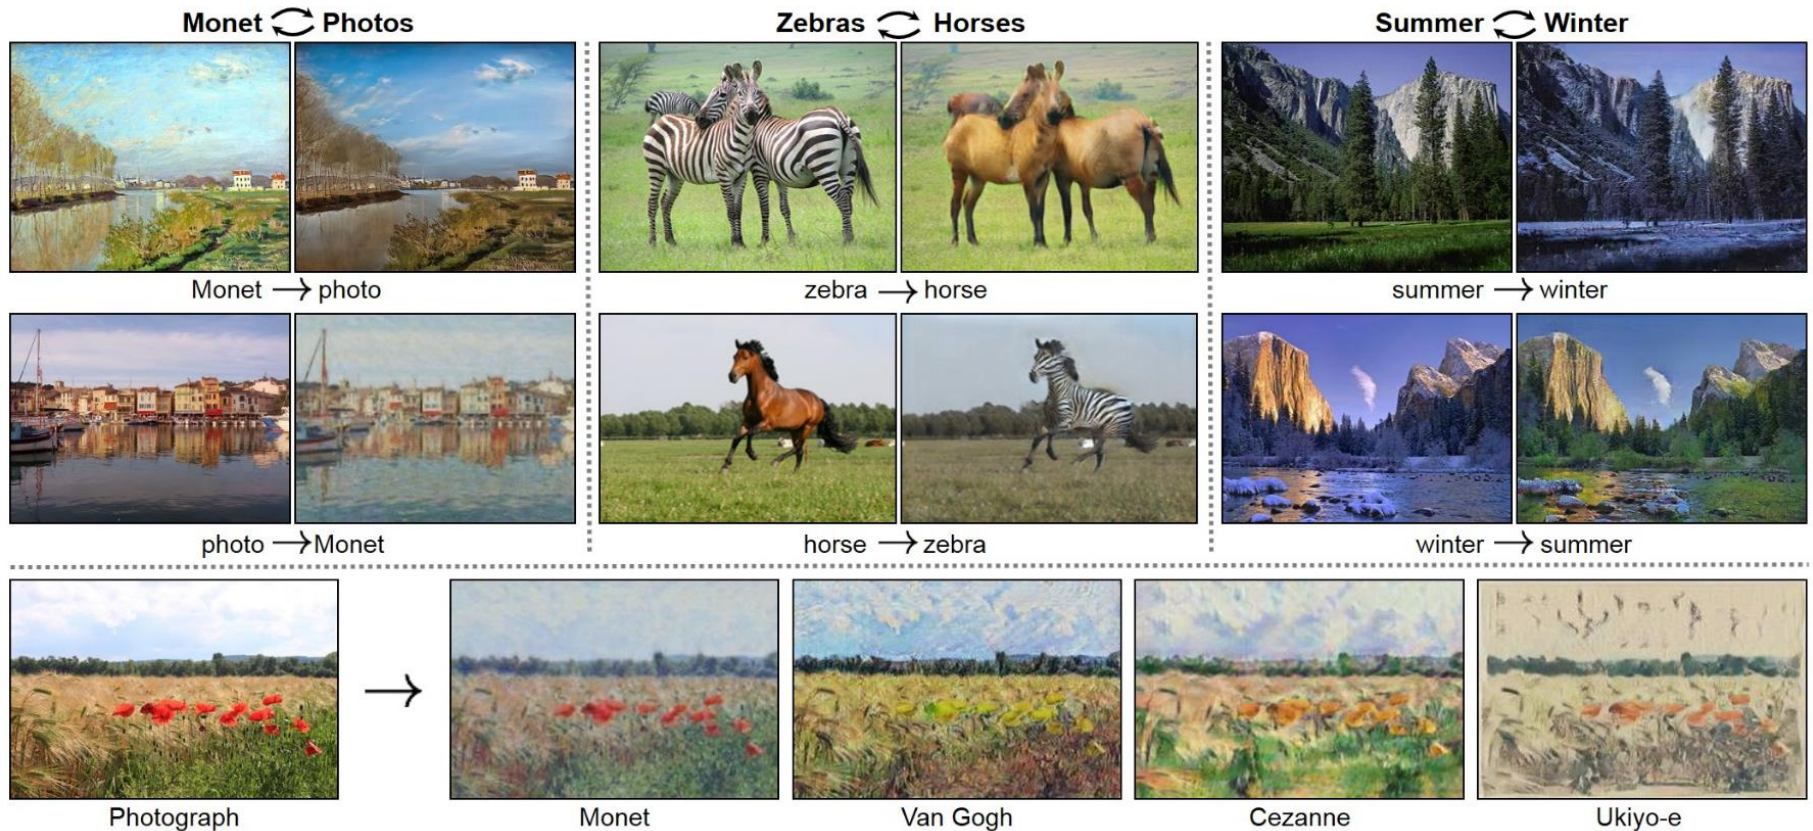


# Application: face generation






Source: <https://www.youtube.com/watch?v=G06dEcZ-QTg>

# Application: style transfer





# Application: image translation





# Conclusion

- GANs provide an attractive way to create a generative model without using an explicit encode (as in variational auto-encoders) but using supervised learning
- However the training is far from trivial and in general finding Nash equilibria in high-dimensional non-convex games is still an important open research problem